

# Stats-506 Problem Set #2

GitHub Repository: <https://github.com/zkl2002/Stats-506/>

Zhekai Liu

## Problem 1 - Modified Random walk

(a).

Step function:

```
#' Sample walk increments
#' @return One time random return with values in {-3,-1,1,10}.
random_step <- function() {
  u1 <- runif(1)
  u2 <- runif(1)
  if (u1 < 0.5) {
    if (u2 < 0.20) -3L else -1L
  } else {
    if (u2 < 0.05) 10L else 1L
  }
}
```

Version 1: using a loop

```
#' Random walk (loop version)
#' @param n Number of steps (non-negative integer).
#' @return Final position of the walk.
random_walk1 <- function(n, seed=NULL) {
  # check input
  if (!(length(n) == 1L && is.numeric(n) && is.finite(n) &&
    n >= 0 && n == as.integer(n))) {
    stop("input must be one non-negative integer")
  }
  if (!is.null(seed)) set.seed(seed)
```

```

n <- as.integer(n)
if (n == 0L) return(0L)
# for loop
pos <- 0L
for (i in seq_len(n)) pos <- pos + random_step()
pos
}

```

## Version 2: using built-in R vectorized function

```

#' Random walk version 2 - built-in R vectorized
#' @param n Number of steps.
#' @return Final position of the walk.
random_walk2 <- function(n, seed=NULL) {
  # check input
  if (!(length(n) == 1L && is.numeric(n) && is.finite(n) &&
        n >= 0 && n == as.integer(n))) {
    stop("input must be one non-negative integer")
  }
  if (!is.null(seed)) set.seed(seed)
  n <- as.integer(n)
  if (n == 0L) return(0L)
  # built-in R vectorized
  u <- runif(2L * n)
  u1 <- u[c(TRUE, FALSE)]
  u2 <- u[c(FALSE, TRUE)]
  dir_pos <- u1 >= 0.5
  inc <- integer(n)
  inc[dir_pos] <- ifelse(u2[dir_pos] < 0.05, 10L, 1L)
  inc[!dir_pos] <- ifelse(u2[!dir_pos] < 0.20, -3L, -1L)
  sum(inc)
}

```

## Version 3: using apply function

```

#' Random walk (inside vapply)
#' @param n Number of steps to take
#' @return Final position after n steps
random_walk3 <- function(n, seed=NULL) {
  # check input
  if (!(length(n) == 1L && is.numeric(n) && is.finite(n) &&

```

```

      n >= 0 && n == as.integer(n))) {
        stop("input must be one non-negative integer")
      }
      # using vapply
      if (!is.null(seed)) set.seed(seed)
      n <- as.integer(n)
      if (n == 0L) return(0L)
      inc <- vapply(seq_len(n), function(i) random_step(), integer(1))
      sum(inc)
    }

```

**Show results:**

```
random_walk1(10)
```

```
[1] 7
```

```
random_walk2(10)
```

```
[1] 2
```

```
random_walk3(10)
```

```
[1] -6
```

```
random_walk1(1000)
```

```
[1] 12
```

```
random_walk2(1000)
```

```
[1] 171
```

```
random_walk3(1000)
```

```
[1] 11
```

**(b).**

```
c(random_walk1(10, seed = 506),
  random_walk2(10, seed = 506),
  random_walk3(10, seed = 506))
```

```
[1] 4 4 4
```

```
c(random_walk1(1000, seed = 506),
  random_walk2(1000, seed = 506),
  random_walk3(1000, seed = 506))
```

```
[1] 73 73 73
```

With same random seed, the result of all three methods are same.

(c).

```
library(microbenchmark)
## n = 1,000
bench_1000 <- microbenchmark(
  loop = random_walk1(1000, seed = 506),
  vec = random_walk2(1000, seed = 506),
  apply = random_walk3(1000, seed = 506)
)
bench_1000
```

Unit: microseconds

expr	min	lq	mean	median	uq	max	neval
loop	2701.6	2788.75	3042.239	2819.85	2912.50	7619.4	100
vec	126.1	150.30	181.712	170.60	200.80	394.8	100
apply	3849.1	3915.00	4234.107	3940.00	4140.75	15110.3	100

```
## n = 100,000
bench_10k <- microbenchmark(
  loop = random_walk1(100000, seed = 506),
  vec = random_walk2(100000, seed = 506),
  apply = random_walk3(100000, seed = 506)
)
bench_10k
```

Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
loop	171.7818	177.1505	193.892952	180.39545	198.4828	308.8743	100	
vec	7.6471	8.1415	9.670794	9.39155	10.2754	17.9583	100	
apply	243.0993	250.8302	269.743757	253.48550	271.6712	464.0754	100	

The fully vectorized implementation is consistently the fastest, then for-loop version comes next, and the apply method approach is the slowest. As the input size grows from 1k to 100k, the performance doesn't change.

(d).

```
reps <- 10000
set.seed(506)
position_10 <- replicate(reps, random_walk2(10, seed = NULL))
cat("The probability that the random walk ends at 0 with 10 steps is:",
    mean(position_10 == 0))
```

The probability that the random walk ends at 0 with 10 steps is: 0.1331

```
set.seed(506)
position_100 <- replicate(reps, random_walk2(100, seed = NULL))
cat("The probability that the random walk ends at 0 with 100 steps is:",
    mean(position_100 == 0))
```

The probability that the random walk ends at 0 with 100 steps is: 0.0184

```
set.seed(506)
position_1000 <- replicate(reps, random_walk2(1000, seed = NULL))
cat("The probability that the random walk ends at 0 with 1000 steps is:",
    mean(position_1000 == 0))
```

The probability that the random walk ends at 0 with 1000 steps is: 0.0049

With Monte Carlo simulation, the probability of random walk ends at 0 is becoming smaller and closer to 0 with steps 10, 100 and 1000.

## Problem 2 - Mean of Mixture of Distributions

```
set.seed(506)
reps <- 10000

daily <- rpois(reps, 8) +
  rnorm(reps, 120, sqrt(24)) +
  rpois(reps, 64) +
  rpois(reps, 72)

mean(daily)
```

```
[1] 264.0015
```

By Monte Carlo simulation, the average number of cars pass this intersection daily is about 264.

## Problem 3 - Linear Regression

```
youtube <- read.csv('https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/youtube/youtube.csv')
```

(a).

```
colnames(youtube)
```

```
[1] "year" "brand"
[3] "superbowl_ads_dot_com_url" "youtube_url"
[5] "funny" "show_product_quickly"
[7] "patriotic" "celebrity"
[9] "danger" "animals"
[11] "use_sex" "id"
[13] "kind" "etag"
[15] "view_count" "like_count"
[17] "dislike_count" "favorite_count"
[19] "comment_count" "published_at"
[21] "title" "description"
[23] "thumbnail" "channel_title"
[25] "category_id"
```

```
drop_cols <- c("brand", "superbowl_ads_dot_com_url", "youtube_url",
              "thumbnail", "channel_title", "published_at",
              "id", "kind", "etag", "title", "description")
youtube_deid <- youtube[ , !(names(youtube) %in% drop_cols)]
youtube <- youtube_deid
dim(youtube)
```


```
[1] 247  14
```

The dimension of the data after de-identify is 247 rows and 14 columns.


**(b).**

**view\_count:**

```
hist(youtube$view_count, main="view_count", xlab="")
```



```
hist(log1p(youtube$view_count), main="log1p(view_count)", xlab="")
```



```
youtube$view_count_log <- log1p(youtube$view_count)
```

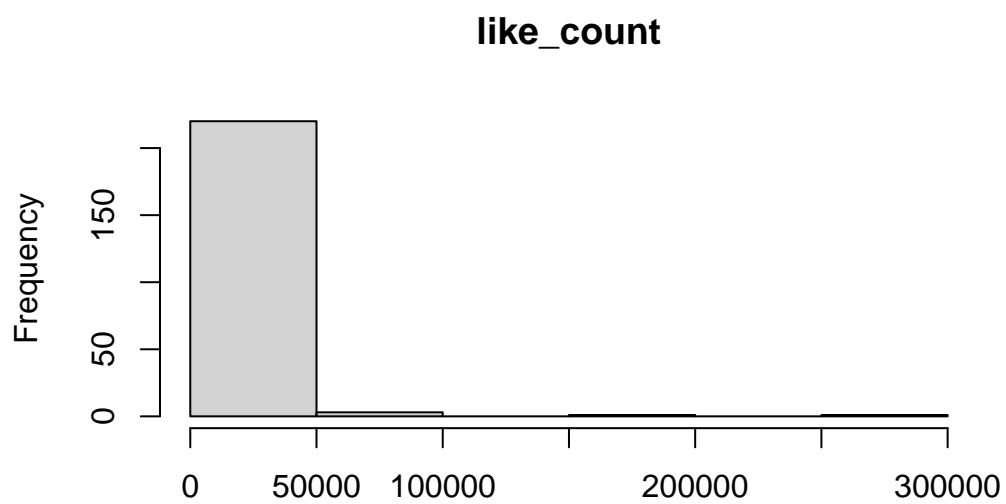
view\_count variable can use a transformation prior to being used as the outcome in a linear regression model.

The view\_count variable is extremely right-skewed, so we apply a  $\log(1 + x)$  transformation on it.

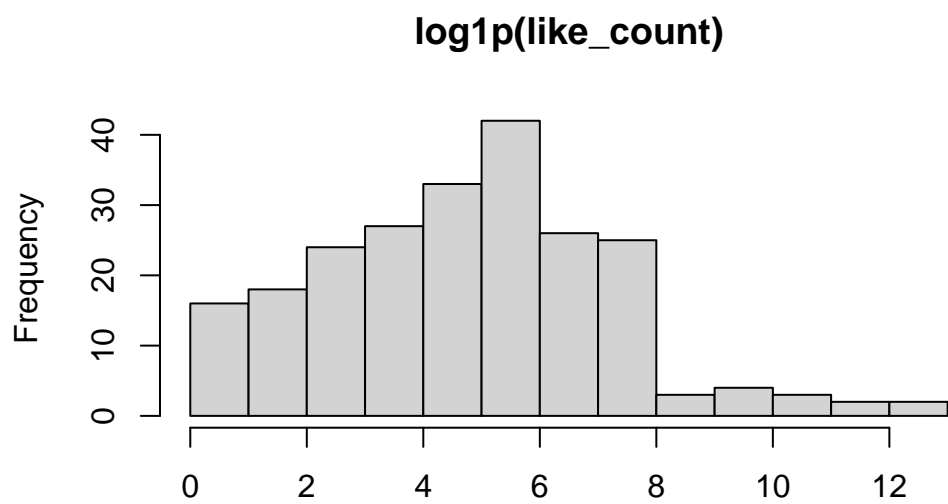
**like\_count:**

```
hist(youtube$like_count, main="like_count", xlab="")
```





```
hist(log1p(youtube$like_count), main="log1p(like_count)", xlab="")
```



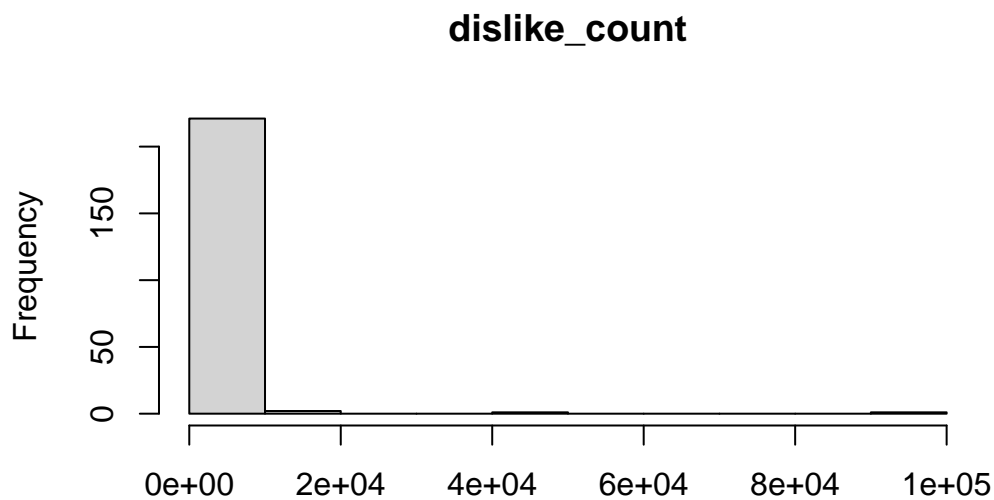
```
youtube$like_count_log <- log1p(youtube$like_count)
```

like\_count variable can use a transformation prior to being used as the outcome in a linear regression model.

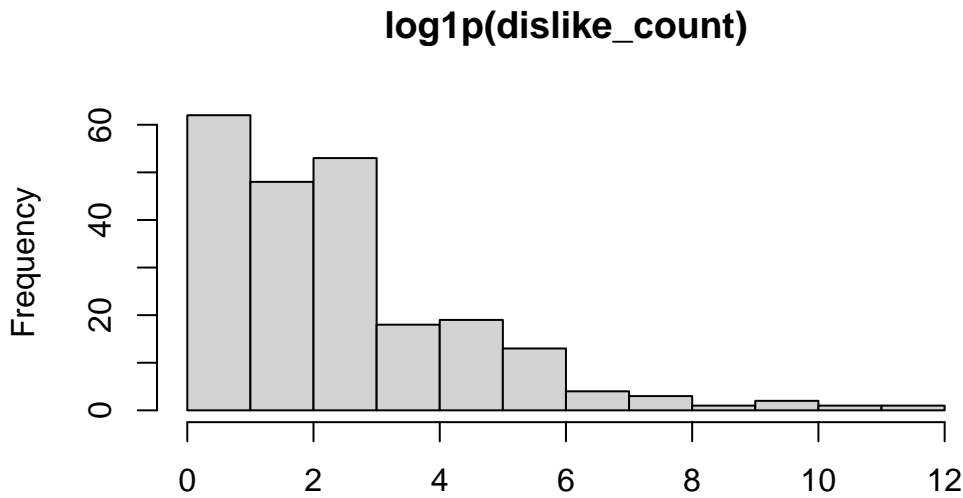
The like\_count variable is extremely right-skewed, so we apply a  $\log(1 + x)$  transformation on it.

**dislike\_count:**

```
hist(youtube$dislike_count, main="dislike_count", xlab="")
```



```
hist(log1p(youtube$dislike_count), main="log1p(dislike_count)", xlab="")
```



```
youtube$dislike_count_log <- log1p(youtube$dislike_count)
```

dislike\_count variable can use a transformation prior to being used as the outcome in a linear regression model.

The dislike\_count variable is extremely right-skewed, so we apply a  $\log(1 + x)$  transformation on it.

**favorite\_count:**


```
unique(youtube$favorite_count)
```

```
[1] 0 NA
```

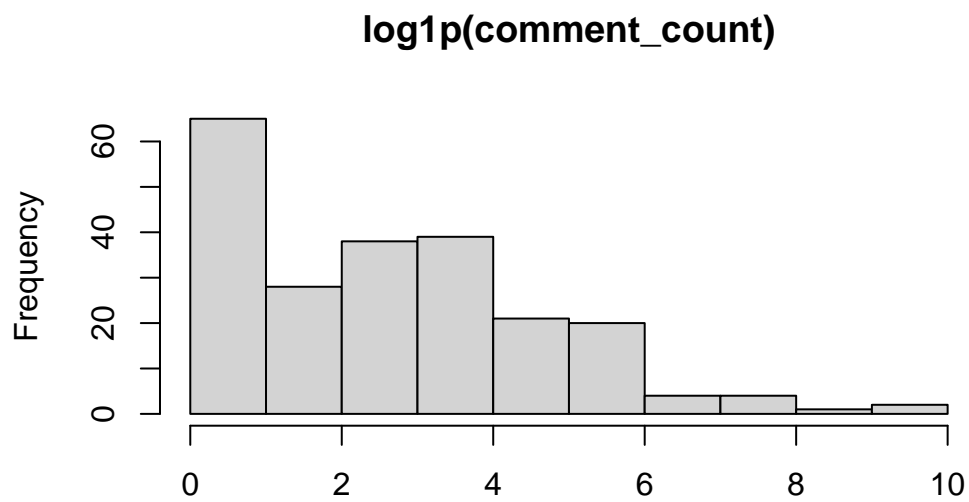
favorite\_count only have 0 and null elements, so this variable would not be appropriate to use as the outcome in a linear regression model.

**comment\_count:**

```
hist(youtube$comment_count, main="comment_count", xlab="")
```



```
hist(log1p(youtube$comment_count), main="log1p(comment_count)", xlab="")
```



```
youtube$comment_count_log <- log1p(youtube$comment_count)
```

comment\_count variable can use a transformation prior to being used as the outcome in a linear regression model.

The comment\_count variable is extremely right-skewed, so we apply a  $\log(1+x)$  transformation on it.

(c).

view\_count:

```
fit_view <- lm(view_count_log ~ year + funny + show_product_quickly +  
               patriotic + celebrity + danger + animals + use_sex,  
               data = youtube)  
summary(fit_view)
```

Call:

```
lm(formula = view_count_log ~ year + funny + show_product_quickly +  
    patriotic + celebrity + danger + animals + use_sex, data = youtube)
```

Residuals:

Min	1Q	Median	3Q	Max
-7.7742	-1.6152	0.1311	1.7036	8.4481

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-31.55016	71.00538	-0.444	0.657
year	0.02053	0.03531	0.582	0.561
funnyTRUE	0.56492	0.46702	1.210	0.228
show_product_quicklyTRUE	0.21089	0.40530	0.520	0.603
patrioticTRUE	0.50699	0.53811	0.942	0.347
celebrityTRUE	0.03548	0.42228	0.084	0.933
dangerTRUE	0.63131	0.41812	1.510	0.132
animalsTRUE	-0.31002	0.39348	-0.788	0.432
use_sexTRUE	-0.38671	0.44782	-0.864	0.389

Residual standard error: 2.787 on 222 degrees of freedom

( 16 )

Multiple R-squared: 0.02694, Adjusted R-squared: -0.008122

F-statistic: 0.7684 on 8 and 222 DF, p-value: 0.631

There's no statistically reliable association between ad attributes and view counts. The estimated directions are all not significant. So we could not provide any significant result.

**like\_count:**

```
fit_like <- lm(like_count_log ~ year + funny + show_product_quickly +
               patriotic + celebrity + danger + animals + use_sex,
               data = youtube)
summary(fit_like)
```

Call:

```
lm(formula = like_count_log ~ year + funny + show_product_quickly +
    patriotic + celebrity + danger + animals + use_sex, data = youtube)
```

Residuals:

Min	1Q	Median	3Q	Max
-5.2860	-1.6333	0.0868	1.4911	7.7431

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-150.51357	63.45723	-2.372	0.0186 *
year	0.07685	0.03155	2.436	0.0157 *
funnyTRUE	0.47476	0.41816	1.135	0.2575
show_product_quicklyTRUE	0.20017	0.36391	0.550	0.5828
patrioticTRUE	0.80689	0.49791	1.621	0.1066
celebrityTRUE	0.41283	0.38212	1.080	0.2812
dangerTRUE	0.63895	0.37350	1.711	0.0886 .
animalsTRUE	-0.05944	0.35298	-0.168	0.8664
use_sexTRUE	-0.42952	0.40064	-1.072	0.2849

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.467 on 216 degrees of freedom

( 22 )

Multiple R-squared: 0.07313, Adjusted R-squared: 0.03881

F-statistic: 2.13 on 8 and 216 DF, p-value: 0.0342

For like\_count, only year shows a positive and statistically significant association with like counts. Danger shows a positive tendency but doesn't reach conventional significance, and the other ad features show no statistically significant associations. As result, only like count numbers would increase with time.

**dislike\_count:**

```
fit_dislike <- lm(dislike_count_log ~ year + funny + show_product_quickly +  
  patriotic + celebrity + danger + animals + use_sex,  
  data = youtube)  
summary(fit_dislike)
```

Call:

```
lm(formula = dislike_count_log ~ year + funny + show_product_quickly +  
  patriotic + celebrity + danger + animals + use_sex, data = youtube)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.0292	-1.3299	-0.3192	0.8986	8.7219

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-183.06813	53.34768	-3.432	0.000719 ***
year	0.09207	0.02653	3.471	0.000626 ***
funnyTRUE	0.25949	0.35154	0.738	0.461224
show_product_quicklyTRUE	0.27511	0.30593	0.899	0.369515
patrioticTRUE	0.81407	0.41859	1.945	0.053095 .
celebrityTRUE	-0.20214	0.32125	-0.629	0.529852
dangerTRUE	0.22184	0.31400	0.707	0.480630
animalsTRUE	-0.21192	0.29675	-0.714	0.475911
use_sexTRUE	-0.32980	0.33681	-0.979	0.328583

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.074 on 216 degrees of freedom

( 22 )

Multiple R-squared: 0.09753, Adjusted R-squared: 0.06411

F-statistic: 2.918 on 8 and 216 DF, p-value: 0.004115

For dislike\_count, only year has a significant positive association with dislike counts. Patriotic shows a positive tendency but doesn't reach conventional significance, and the other ad features show no statistically significant associations. These also means with time increase, dislike count would also increase.

**comment\_count:**

```
fit_comment <- lm(comment_count_log ~ year + funny + show_product_quickly +
  patriotic + celebrity + danger + animals + use_sex,
  data = youtube)
summary(fit_comment)
```

Call:

```
lm(formula = comment_count_log ~ year + funny + show_product_quickly +
  patriotic + celebrity + danger + animals + use_sex, data = youtube)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.1372	-1.4665	-0.1427	1.4061	5.8468

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-99.09835	52.92351	-1.872	0.0625 .
year	0.05034	0.02632	1.913	0.0571 .
funnyTRUE	0.21954	0.34528	0.636	0.5256
show_product_quicklyTRUE	0.40939	0.30229	1.354	0.1771
patrioticTRUE	0.66698	0.39902	1.672	0.0961 .
celebrityTRUE	0.29767	0.31541	0.944	0.3464
dangerTRUE	0.17784	0.31069	0.572	0.5677
animalsTRUE	-0.26802	0.29347	-0.913	0.3621
use_sexTRUE	-0.39323	0.33163	-1.186	0.2370

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.039 on 213 degrees of freedom

( 25 )

Multiple R-squared: 0.06535, Adjusted R-squared: 0.03025

F-statistic: 1.862 on 8 and 213 DF, p-value: 0.06748

For log comment counts, there are only marginal positive associations with year and patriotic, but they do not reach the conventional 0.05 significance level. The other ad features show no statistically reliable associations, and the model's explanatory power is low.

(d).



```

vars <- c("view_count_log","year","funny","show_product_quickly",
          "patriotic","celebrity","danger","animals","use_sex")
dat <- na.omit(youtube[, vars])

X <- model.matrix(view_count_log ~ year + funny + show_product_quickly +
                  patriotic + celebrity + danger + animals + use_sex,
                  data = dat)
y <- dat$view_count_log
beta_hat <- solve(t(X) %*% X, t(X) %*% y)
beta_hat

```

	[,1]
(Intercept)	-31.55015804
year	0.02053399
funnyTRUE	0.56492445
show_product_quicklyTRUE	0.21088918
patrioticTRUE	0.50699051
celebrityTRUE	0.03547862
dangerTRUE	0.63131085
animalsTRUE	-0.31001838
use_sexTRUE	-0.38670726

The result is same with the lm model in part c.